

ChatGPT를 API에서 이용하기

3-3 입력과 출력의 길이 제한 및 요금에 영향을 주는 '토큰'

Tokenizer와 tiktoken 소개

```
In [1]: !pip install tiktoken
```

```
Requirement already satisfied: tiktoken in c:\programdata\anaconda3\lib\site-packages (0.12.0)
Requirement already satisfied: regex>=2022.1.18 in c:\programdata\anaconda3\lib\site-packages (from tiktoken) (2024.11.6)
Requirement already satisfied: requests>=2.26.0 in c:\programdata\anaconda3\lib\site-packages (from tiktoken) (2.32.5)
Requirement already satisfied: charset_normalizer<4,>=2 in c:\programdata\anaconda3\lib\site-packages (from requests>=2.26.0->tiktoken) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in c:\programdata\anaconda3\lib\site-packages (from requests>=2.26.0->tiktoken) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\programdata\anaconda3\lib\site-packages (from requests>=2.26.0->tiktoken) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in c:\programdata\anaconda3\lib\site-packages (from requests>=2.26.0->tiktoken) (2025.10.5)
```

```
In [2]: import tiktoken
```

```
text = "It's easy to make something cool with LLMs, but very hard to make something production-ready with them."
encoding = tiktoken.encoding_for_model("gpt-4.1-mini")
tokens = encoding.encode(text)
print(len(tokens))
```

23

```
In [3]: text = "LLM을 사용해서 먼저 보이는 것을 만들기는 쉽지만, 프로덕션 수준으로 만들어 내기는 매우 어렵다."
```

```
encoding = tiktoken.encoding_for_model("gpt-4.1-mini")
tokens = encoding.encode(text)
print(len(tokens))
```

3-4 Chat Completions API를 접하는 환경 준비

OpenAI API 키 준비

```
In [4]: import os  
from dotenv import load_dotenv  
load_dotenv()  
OPENAI_API_KEY = os.getenv("OPENAI_API_KEY")
```

3-5 Chat Completions API 사용해 보기

OpenAI 라이브러리

```
In [5]: !pip install openai
```

```
Requirement already satisfied: openai in c:\programdata\anaconda3\lib\site-packages (2.15.0)
Requirement already satisfied: aiohttp<2.0.0,>=1.3.4 in c:\programdata\anaconda3\lib\site-packages (from openai) (4.7.0)
Requirement already satisfied: certifi in c:\programdata\anaconda3\lib\site-packages (from openai) (2025.10.5)
Requirement already satisfied: idna>=2.8 in c:\programdata\anaconda3\lib\site-packages (from aiohttp<2.0.0,>=1.3.4->openai) (3.7)
Requirement already satisfied: httpcore==1.* in c:\programdata\anaconda3\lib\site-packages (from aiohttp<2.0.0,>=1.3.4->openai) (1.0.9)
Requirement already satisfied: h11>=0.16 in c:\programdata\anaconda3\lib\site-packages (from httpcore==1.*->aiohttp<2.0.0,>=1.3.4->openai) (0.16.0)
Requirement already satisfied: annotated-types>=0.6.0 in c:\programdata\anaconda3\lib\site-packages (from pydantic<3,>=1.9.0->openai) (0.6.0)
Requirement already satisfied: pydantic-core==2.27.1 in c:\programdata\anaconda3\lib\site-packages (from pydantic<3,>=1.9.0->openai) (2.27.1)
Requirement already satisfied: colorama in c:\programdata\anaconda3\lib\site-packages (from tqdm>4->openai) (0.4.6)
```

Chat Completions API 호출

```
In [6]: from openai import OpenAI

client = OpenAI()

response = client.chat.completions.create(
    model="gpt-4.1-mini",
    messages=[
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "Hello! I'm John."}
    ]
)

print(response)
```

```
ChatCompletion(id='chatcmpl-CxlgfESqM8uM7Tspp9ZxbYCwRujW1', choices=[Choice(finish_reason='stop', index=0, logprobs=None, message=ChatCompletionMessage(content='Hello John! How can I assist you today?', refusal=None, role='assistant', annotations=[], audio=None, function_call=None, tool_calls=None)), created=1768360697, model='gpt-4.1-mini-2025-04-14', object='chat.completion', service_tier='default', system_fingerprint='fp_376a7cce1', usage=CompletionUsage(completion_tokens=10, prompt_tokens=22, total_tokens=32, completion_tokens_details=CompletionTokensDetails(accepted_prediction_tokens=0, audio_tokens=0, reasoning_tokens=0, rejected_prediction_tokens=0), prompt_tokens_details=PromptTokensDetails(audio_tokens=0, cached_tokens=0)))
```

```
In [7]: print(response.model_dump_json(indent=2))
```

```
{  
    "id": "chatcmpl-CxlgfESqM8uM7Tspp9ZxbYCwRujW1",  
    "choices": [  
        {  
            "finish_reason": "stop",  
            "index": 0,  
            "logprobs": null,  
            "message": {  
                "content": "Hello John! How can I assist you today?",  
                "refusal": null,  
                "role": "assistant",  
                "annotations": [],  
                "audio": null,  
                "function_call": null,  
                "tool_calls": null  
            }  
        }  
    ],  
    "created": 1768360697,  
    "model": "gpt-4.1-mini-2025-04-14",  
    "object": "chat.completion",  
    "service_tier": "default",  
    "system_fingerprint": "fp_376a7cce1",  
    "usage": {  
        "completion_tokens": 10,  
        "prompt_tokens": 22,  
        "total_tokens": 32,  
        "completion_tokens_details": {  
            "accepted_prediction_tokens": 0,  
            "audio_tokens": 0,  
            "reasoning_tokens": 0,  
            "rejected_prediction_tokens": 0  
        },  
        "prompt_tokens_details": {  
            "audio_tokens": 0,  
            "cached_tokens": 0  
        }  
    }  
}
```

대화 능력을 바탕으로 한 응답 얻기

```
In [8]: response = client.chat.completions.create(  
    model="gpt-4.1-mini",  
    messages=[  
        {"role": "system", "content": "You are a helpful assistant."},  
        {"role": "user", "content": "Hello! I'm John."},  
        {"role": "assistant", "content": "Hello John! How can I assist you today?"},  
        {"role": "user", "content": "Do you know my name?"}  
    ]  
)  
  
print(response.model_dump_json(indent=2))
```

```
{  
    "id": "chatcmpl-CxlggiklDhumMCU7arBjT90RVm1Wk",  
    "choices": [  
        {  
            "finish_reason": "stop",  
            "index": 0,  
            "logprobs": null,  
            "message": {  
                "content": "Yes, you mentioned that your name is John. How can I help you today, John?",  
                "refusal": null,  
                "role": "assistant",  
                "annotations": [],  
                "audio": null,  
                "function_call": null,  
                "tool_calls": null  
            }  
        }  
    ],  
    "created": 1768360698,  
    "model": "gpt-4.1-mini-2025-04-14",  
    "object": "chat.completion",  
    "service_tier": "default",  
    "system_fingerprint": "fp_376a7cce1",  
    "usage": {  
        "completion_tokens": 19,  
        "prompt_tokens": 46,  
        "total_tokens": 65,  
        "completion_tokens_details": {  
            "accepted_prediction_tokens": 0,  
            "audio_tokens": 0,  
            "reasoning_tokens": 0,  
            "rejected_prediction_tokens": 0  
        },  
        "prompt_tokens_details": {  
            "audio_tokens": 0,  
            "cached_tokens": 0  
        }  
    }  
}
```

응답을 스트리밍으로 받기

```
In [9]: response = client.chat.completions.create(
    model="gpt-4.1-mini",
    messages=[
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "Hello! I'm John."}
    ],
    stream=True
)

for chunk in response:
    choice = chunk.choices[0]
    if choice.finish_reason is None:
        print(choice.delta.content)
```

```
Hello
John
!
How
can
I
assist
you
today
?
```

(칼럼) Completions API

```
In [10]: from openai import OpenAI

client = OpenAI()

response = client.chat.completions.create(
    model="gpt-4.1-mini",
    messages=[
        {"role": "user", "content": "Hello! I'm John."}
    ]
```

```
)  
print(response.model_dump_json(indent=2))
```

```
{  
    "id": "chatcmpl-CxlghMjgMEFIsoVbXLPhZsrbcIVjk",  
    "choices": [  
        {  
            "finish_reason": "stop",  
            "index": 0,  
            "logprobs": null,  
            "message": {  
                "content": "Hello John! How can I assist you today?",  
                "refusal": null,  
                "role": "assistant",  
                "annotations": [],  
                "audio": null,  
                "function_call": null,  
                "tool_calls": null  
            }  
        }  
    ],  
    "created": 1768360699,  
    "model": "gpt-4.1-mini-2025-04-14",  
    "object": "chat.completion",  
    "service_tier": "default",  
    "system_fingerprint": "fp_376a7cce1",  
    "usage": {  
        "completion_tokens": 10,  
        "prompt_tokens": 12,  
        "total_tokens": 22,  
        "completion_tokens_details": {  
            "accepted_prediction_tokens": 0,  
            "audio_tokens": 0,  
            "reasoning_tokens": 0,  
            "rejected_prediction_tokens": 0  
        },  
        "prompt_tokens_details": {  
            "audio_tokens": 0,  
            "cached_tokens": 0  
        }  
    }  
}
```

```
In [11]: messages = [
    {"role": "user", "content": "Hello! I'm John."},
    {"role": "assistant", "content": "Nice to meet you, John!"},
    {"role": "user", "content": "Do you know my name?"}
]

response = client.chat.completions.create(
    model="gpt-4.1-mini",
    messages=messages
)
print(response.model_dump_json(indent=2))
```

```
{  
    "id": "chatcmpl-CxlgisBESB3VNHOBDm3L0xJywyci0",  
    "choices": [  
        {  
            "finish_reason": "stop",  
            "index": 0,  
            "logprobs": null,  
            "message": {  
                "content": "Yes, you mentioned that your name is John. How can I assist you today?",  
                "refusal": null,  
                "role": "assistant",  
                "annotations": [],  
                "audio": null,  
                "function_call": null,  
                "tool_calls": null  
            }  
        }  
    ],  
    "created": 1768360700,  
    "model": "gpt-4.1-mini-2025-04-14",  
    "object": "chat.completion",  
    "service_tier": "default",  
    "system_fingerprint": "fp_376a7cce1",  
    "usage": {  
        "completion_tokens": 17,  
        "prompt_tokens": 33,  
        "total_tokens": 50,  
        "completion_tokens_details": {  
            "accepted_prediction_tokens": 0,  
            "audio_tokens": 0,  
            "reasoning_tokens": 0,  
            "rejected_prediction_tokens": 0  
        },  
        "prompt_tokens_details": {  
            "audio_tokens": 0,  
            "cached_tokens": 0  
        }  
    }  
}
```

3-6 Function calling

Function calling의 샘플 코드

OpenAI 공식 문서를 바탕으로 일부 수정한 코드입니다.

```
In [12]: import json

def get_current_weather(location, unit="celsius"):
    weather_info = {
        "location": location,
        "temperature": "25",
        "unit": "celsius",
        "forecast": ["sunny", "windy"],
    }
    return json.dumps(weather_info)
```

```
In [13]: functions = [
    {
        "name": "get_current_weather",
        "description": "Get the current weather in a given location",
        "parameters": {
            "type": "object",
            "properties": {
                "location": {
                    "type": "string",
                    "description": "The city and state, e.g. Seoul",
                },
                "unit": {"type": "string", "enum": ["celsius", "fahrenheitz"]},
            },
            "required": ["location"]
        }
    }
]
```

```
In [14]: messages = [{"role": "user", "content": "What's the weather like in Seoul?"}]
```

```
response = client.chat.completions.create(
    model="gpt-4.1-mini",
    messages=messages,
    functions=functions
)

print(response.model_dump_json(indent=2))
```

```
{  
    "id": "chatcmpl-CxlgjyvZfVduUJuWCRdPJo20vAsEG",  
    "choices": [  
        {  
            "finish_reason": "function_call",  
            "index": 0,  
            "logprobs": null,  
            "message": {  
                "content": null,  
                "refusal": null,  
                "role": "assistant",  
                "annotations": [],  
                "audio": null,  
                "function_call": {  
                    "arguments": "{\"location\": \"Seoul\", \"unit\": \"celsius\"}",  
                    "name": "get_current_weather"  
                },  
                "tool_calls": null  
            }  
        }  
    ],  
    "created": 1768360701,  
    "model": "gpt-4.1-mini-2025-04-14",  
    "object": "chat.completion",  
    "service_tier": "default",  
    "system_fingerprint": "fp_376a7cce1",  
    "usage": {  
        "completion_tokens": 21,  
        "prompt_tokens": 76,  
        "total_tokens": 97,  
        "completion_tokens_details": {  
            "accepted_prediction_tokens": 0,  
            "audio_tokens": 0,  
            "reasoning_tokens": 0,  
            "rejected_prediction_tokens": 0  
        },  
        "prompt_tokens_details": {  
            "audio_tokens": 0,  
            "cached_tokens": 0  
        }  
    }  
}
```

```
        }
    }

In [15]: response_message = response.choices[0].message

available_functions = {
    "get_current_weather": get_current_weather,
}
function_name = response_message.function_call.name
function_to_call = available_functions[function_name]
function_args = json.loads(response_message.function_call.arguments)

function_response = function_to_call(
    location=function_args.get("location"),
    unit=function_args.get("unit"),
)

print(function_response)
```

```
{"location": "Seoul", "temperature": "25", "unit": "celsius", "forecast": ["sunny", "windy"]}
```

```
In [16]: messages.append(response_message)
messages.append(
{
    "role": "function",
    "name": function_name,
    "content": function_response,
}
)
```

```
In [17]: print(messages)
```

```
[{'role': 'user', 'content': "What's the weather like in Seoul?"}, ChatCompletionMessage(content=None, refusal=None, role='assistant', annotations=[], audio=None, function_call=FunctionCall(arguments='{"location": "Seoul", "unit": "celsius"}', name='get_current_weather'), tool_calls=None), {'role': 'function', 'name': 'get_current_weather', 'content': '{"location": "Seoul", "temperature": "25", "unit": "celsius", "forecast": ["sunny", "windy"]}'}]
```

```
In [18]: second_response = client.chat.completions.create(
    model="gpt-4.1-mini",
    messages=messages,
)
```

```
print(second_response.model_dump_json(indent=2))
```

```
{  
    "id": "chatcmpl-CxlgjqxPcjfPwUPTNIbdzdmeB7nge",  
    "choices": [  
        {  
            "finish_reason": "stop",  
            "index": 0,  
            "logprobs": null,  
            "message": {  
                "content": "The weather in Seoul is currently sunny and windy, with a temperature of 25°C.",  
                "refusal": null,  
                "role": "assistant",  
                "annotations": [],  
                "audio": null,  
                "function_call": null,  
                "tool_calls": null  
            }  
        }  
    ],  
    "created": 1768360701,  
    "model": "gpt-4.1-mini-2025-04-14",  
    "object": "chat.completion",  
    "service_tier": "default",  
    "system_fingerprint": "fp_376a7cce1",  
    "usage": {  
        "completion_tokens": 18,  
        "prompt_tokens": 76,  
        "total_tokens": 94,  
        "completion_tokens_details": {  
            "accepted_prediction_tokens": 0,  
            "audio_tokens": 0,  
            "reasoning_tokens": 0,  
            "rejected_prediction_tokens": 0  
        },  
        "prompt_tokens_details": {  
            "audio_tokens": 0,  
            "cached_tokens": 0  
        }  
    }  
}
```

